

## **1. Problem Being Solved**

Modern companies communicate with customers across:

- web forms
- emails
- WhatsApp / chat
- multiple meetings

Today, this context is:

- fragmented across tools
- lost between meetings
- re-asked repeatedly
- poorly transferred between Sales → Engineering → PM

### **Result:**

Missed requirements, scope creep, poor proposals, repeated clarification meetings.

---

## **2. Vision of the Proposed System**

This architecture designs an **AI-first, context-aware enterprise communication system** where:

- AI participates **before, during, and after** meetings
- AI remembers **everything that matters**, across all channels
- AI proactively prepares humans (instead of waiting for prompts)
- Context improves over time — the system learns, not just the model

The system behaves like a **continuously learning project assistant**, not a chatbot.

---

## **3. Core Architectural Principle**

### **Memory > Model**

The intelligence comes from:

- how context is stored

- how it is retrieved
- how it is versioned
- how it is validated

LLMs reason over **carefully curated context packs**, not raw history.

---

## 1. High-Level System Blocks

**Channels → Ingestion → Memory Platform → AI Intelligence → Outputs → Humans → back to Channels**

### Key Blocks

#### 1. Channels

- Web / Forms
- Email
- WhatsApp / Chat
- Meetings (Zoom / Teams / Meet)

#### 2. Ingestion Gateway

- Normalizes all inputs
- Tags metadata
- Routes events to storage and workflows

#### 3. Context & Memory Platform

- Structured DB (state)
- Vector DB (semantic memory)
- Memory Graph (decisions, dependencies, risks)
- Artifact Store (raw files)

#### 4. AI Intelligence Services

- Scheduling
- Pre-Meeting Intelligence

- Live Meeting Bot
- Post-Meeting Update
- Engineering, Scope & Proposal generation

## 5. Outputs + Humans

- Internal dashboard
  - CRM
  - Proposal documents
- 

## 2. Why Multiple Memory Types Are Needed

### Memory Type Purpose

Structured DB What is open/closed, owners, status

Vector DB “Find where this was discussed before”

Memory Graph Decisions → risks → dependencies

Artifact Store Source of truth (audio, PDFs, emails)

Context Ledger What changed, when, and why

This separation prevents hallucination and enables explainability.

---

## STEP 1 — Customer Request Intake

### What happens

- Customer sends request via web/email/WhatsApp
- System extracts intent, requirements, urgency
- Sales receives an AI-generated internal brief

### Tech Stack

- Web: Next.js / React
- Email: Gmail API / Microsoft Graph

- WhatsApp: WhatsApp Business API / Twilio
- LLM extraction: GPT-4.1 / GPT-4o / GPT-5
- Storage: Postgres + Vector DB + Object Store

#### Why

- Ensures every request enters the same pipeline
  - Prevents context loss at the first interaction
- 

## STEP 2 — External Enrichment

#### What happens

- System fetches customer website
- AI summarizes industry, maturity, tech hints
- Marks all enrichment as *inferred*, not confirmed

#### Tech Stack

- Website extraction: Diffbot / Mercury Parser
- Summarization: GPT-4.1 / GPT-4o
- Embeddings: OpenAI text-embedding-3-\*
- Storage: Vector DB + Knowledge Graph

#### Why

- Sales and Engineering start meetings informed
  - Reduces basic “who are you?” questions
- 

## STEP 3 — Sales Notification & Meeting Draft

#### What happens

- AI creates a short Sales Brief
- CRM opportunity is created
- Meeting intent is defined (discovery/follow-up)

## Tech Stack

- CRM: Salesforce / HubSpot
- Notifications: Slack API / Teams API
- LLM: GPT-4.1 (summarization only)

## Why

- AI assists humans, does not replace decision-making
- 

## STEP 4 — AI-Based Scheduling

### What happens

- Finds meeting times using:
  - availability
  - time zones
  - **personal preferences**
- Suggests optimal slots to customer

## Tech Stack

- Calendar APIs: Google Calendar / Microsoft Graph
- Constraint solving: OR-Tools
- LLM (language only): GPT-4o

## Why

- Calendars show availability, not preference
  - Scheduling is deterministic, not hallucinated
- 

## STEP 5 — Pre-Meeting Context Intelligence

### What happens

- Runs at T-24h (deep brief) and T-2h (delta brief)
- Retrieves all relevant history

- Detects:
  - gaps
  - contradictions
  - risks
- Produces role-specific questions + agenda

## Tech Stack

- Vector DB: Pinecone / Weaviate / Milvus
- Graph DB: Neo4j / Neptune
- Orchestration: LangGraph + Temporal
- LLM reasoning (RAG): GPT-5 / GPT-4.1

## Why

- This is the system's **core intelligence**
  - Humans walk into meetings prepared
- 

## PAGE 4 — Workflow Steps 6–10 (During & After Meetings)

### STEP 6 — Live Meeting Bot (During Meeting)

#### What happens

- AI joins meeting
- Records audio
- Transcribes speech
- Identifies speakers and roles

## Tech Stack

- Meeting APIs: Zoom / Teams
- STT: Whisper / Deepgram / AssemblyAI
- Diarization: pyannote
- Optional products: Otter.ai, Fireflies, Gong

## Why

- Meetings are the highest-value data source
  - Accuracy here affects everything downstream
- 

## STEP 7 — Post-Meeting Update

### What happens

- AI summarizes meeting
- Extracts decisions, answers, open questions
- Updates memory with versioned changes

### Tech Stack

- LLM extraction: GPT-4.1 / GPT-5
- Validation: JSON schema + rule checks
- Storage: DB + Graph + Ledger

## Why

- Converts conversation into structured truth
  - Prevents repeated discussions in future meetings
- 

## STEP 8 — Engineering Context Usage

### What happens

- Engineers receive a clean technical brief
- AI suggests architecture options + assumptions

### Tech Stack

- RAG pipeline: LlamaIndex / LangChain
- LLM: GPT-4.1 / GPT-5
- Knowledge sources: internal design docs

## Why

- Engineers start from context, not guesswork
- 

## STEP 9 — Technical Questionnaire

### What happens

- AI detects missing technical details
- Generates prioritized clarification questions

### Tech Stack

- LLM: GPT-4.1 / GPT-5
- Checklists + Graph queries

### Why

- Prevents late-stage surprises
- 

## STEP 10 — Follow-Up Meetings

### What happens

- Steps 5 → 6 → 7 repeat
- Only unresolved or changed items are focused

### Why

- Context accumulates
  - Meetings become shorter and sharper
- 

## PAGE 5 — Finalization, Governance & Market Reality

## STEP 11 — Scope Definition

### What happens

- AI drafts in-scope / out-of-scope
- Humans approve and lock baseline

### Tech Stack

- LLM: GPT-4.1 / GPT-5
  - Graph validation
- 

## **STEP 12 — Project Planning**

### **What happens**

- Timeline, milestones, roles generated
- PM validates against real capacity

### **Tech Stack**

- LLM: GPT-4.1
  - Optional integration: Jira / Azure DevOps
- 

## **STEP 13 — Proposal Generation**

### **What happens**

- AI fills proposal template
- No pricing invented
- Human approval required

### **Tech Stack**

- LLM: GPT-4.1 / GPT-5
  - Templates: Google Docs API / docxtpl
  - Export: PDF/DOCX tools
- 

## **STEP 14 — Continuous Context Loop**

### **What happens**

- New emails/chats restart the loop
- AI checks if context changes
- Triggers new briefs if needed

---

## Governance, Safety & Reality Check

- RAG is mandatory for reasoning steps
  - Fine-tuning is optional and future-stage
  - Humans approve customer-facing outputs
  - Context ledger enables audits
- 

## Market Products vs This Architecture

Product	Coverage
Otter / Fireflies	Meeting capture only
Gong / Chorus	Sales conversations
Notion AI	Document assistance
Salesforce Einstein CRM	CRM insights

**None** provide:

- cross-channel memory
- decision graphs
- proactive pre-meeting intelligence

That's the unique value of this system.